

OS 2020 Spring Project1 Report

系級:資工二 學號: b07902064 姓名:蔡銘軒

1. Design

- Common design for all scheduling algorithms:
 - A process is defined by `struct Process`. The member variables are:
 - name: the name of the process
 - pid: the actual process id assigned by the kernel.
 - ready time/execution time: Information given in the input. Critical for scheduling decision.
 - remaining time: the remaining execution time of the process. Used in RR and PSJF.
 - index: the index of the process in the sorted array according to ready time. Used for quick access and making process management convenient.
 - After reading the input into an array of `struct Process`, I sort the processes by their ready time. This is needed regardless of the choice of the scheduling algorithms.
 - Run a for loop simulating the actual time flow. Using the variable `current_time` to indicate the time passed from the beginning.
 - When I set a process running, I set its priority to the highest possible value (in this case, 99), so the real scheduler in the kernel would execute the process.
 - When process A *preempts* process B, I set the priority of process A to the lowest possible value (in this case, 1), and set process B running (by setting its priority to the highest value). This way, the real scheduler in the kernel would execute process B instead of process A (hopefully).

- FIFO:

Since the array of processes is sorted, at every time unit, I simply check two things:

1. If there are no process running at this point, check if the next process has arrived (`ready_time ≤ current_time`). If true, run the process.
2. If there is a running process, check if it has finished its execution. If not, continue the loop; otherwise, remove the process so the next process can be scheduled.

- RR:

In the Round-Robin algorithm, I use a queue to maintain the flow of the processes. The queue is defined in the `queue.c` and `queue.h` files. In each time unit, I check the following things:

1. Check if there are any processes ready at this moment. If yes, append all of them to the queue.
2. If there is a running process, check if it has used up its time slice (quantum). If yes, check if it has finished its execution. If yes, simply remove the process; otherwise, preempt the current process with the next process in line. After setting the being preempted, move the process to the end of the queue to wait for the next execution.
3. If there is no process running and the queue is not empty, run the process at the beginning of the queue and refresh the time slice (quantum) for the newly run process.

- SJF:

In the Shortest Job First algorithm, I use a priority queue (min heap) to manage the processes. The heap is defined in the `heap.c` and `heap.h` files. In each time unit, I check the following things:

1. Check if there are any processes ready at this moment. If yes, insert all of them into the heap.
2. If there is process waiting in the heap, check if any process is currently running. If yes, continue (SJF is non-preemptive, we cannot preempt the running process); otherwise, run the process with the shortest execution time in the heap and remove it from the heap.
3. If a process is running, check if it has finished execution. If yes, remove the process.

■ PSJF:

In the Preemptive Shortest Job First algorithm, I use a priority queue identical to the one used in SJF. In each time unit, I check the following things:

1. If there is a process running and it has finished execution, remove the process from execution and from the heap.
2. If there are any processes ready at this moment, insert all of them to the heap.
3. If the heap is non-empty, check if there is any process running. If not, run the process in the heap that has the shortest remaining time; otherwise, compare the remaining time of the current process and the one in the heap. If the current process has less remaining time, continue; otherwise preempt the current process with the one in the heap.

2. Kernel Version

In this project, the kernel version is 4.14.25 x86_64. And the operating system is Ubuntu 16.04 server.

3. Performance Analysis

In the following chart, the numbers indicate how much longer a process takes to finish execution on average in this project compared to theoretical scheduling result.

	Error
FIFO_1	1.26%
FIFO_2	1.00%
FIFO_3	-0.20%
FIFO_4	0.63%
FIFO_5	0.56%
PSJF_1	0.56%
PSJF_2	1.26%
PSJF_3	1.25%
PSJF_4	0.99%
PSJF_5	5.83%
RR_1	3.39%

RR_2	0.82%
RR_3	0.55%
RR_4	-0.37%
RR_5	0.79%
SJF_1	0.38%
SJF_2	3.19%
SJF_3	2.20%
SJF_4	1.28%
SJF_5	1.02%

The complete result returned by my program is as follows (the program is not included in the project. I used it for self evaluation only.):

```
FIFO_1_stdout.txt 0.012697638236726042
FIFO_2_stdout.txt 0.010032171733699748
FIFO_3_stdout.txt -0.0020142972851228264
FIFO_4_stdout.txt 0.006359720787742643
FIFO_5_stdout.txt 0.005693467204774341
PSJF_1_stdout.txt 0.005667820544554644
PSJF_2_stdout.txt 0.012652231941444356
PSJF_3_stdout.txt 0.01256360840556093
PSJF_4_stdout.txt 0.009981987462084155
PSJF_5_stdout.txt 0.0583777050260505
RR_1_stdout.txt 0.03391037427588189
RR_2_stdout.txt 0.008298850181209086
RR_3_stdout.txt 0.0055025037992447205
RR_4_stdout.txt -0.0037277886760344416
RR_5_stdout.txt 0.007944871250410924
SJF_1_stdout.txt 0.0038785918271535052
SJF_2_stdout.txt 0.03191950591115621
SJF_3_stdout.txt 0.022017364567404805
SJF_4_stdout.txt 0.012828081685027063
SJF_5_stdout.txt 0.010201213770785972
```

In most of the cases, the processes took longer to execute than the theoretical expectation. I think this is reasonable since my scheduler is, after all, a user process that is scheduled by the actual scheduler in the kernel. The execution of child processes are also at the mercy of the real scheduler. There were some processes irrelevant to this project but still needed to be run by the kernel. Also, the scheduler in this project involves forking and managing child processes, which add more complexity to my scheduler than the theoretical implementation, which led to longer execution time.

In some cases, however, some processes finished execution earlier than my scheduler expected them to. A possible reason is that, when a process was supposed to be waiting in the queue, it was actually run by the kernel scheduler. Or if we consider the moment when a process is scheduled to be run next, but hasn't actually been run (the moment when its priority should be set to 99, but just before that operation), every process has the same priority, 1. At this moment, the scheduler in kernel may choose any of these processes to run, despite the fact that they should be waiting in the queue.